



קורס מערכות לומדות קיץ 2021

הערות ודגשים לגבי התרגילים הרטובים

סיכום נושאים שהועברו

הכנה **לקרב**



קורס מערכות לומדות קיץ 2021

הערות ודגשים לגבי התרגילים הרטובים

סיכום נושאים שהועברו

הכנה להאקט'ון

Session 2




- Ex. 0: basic script importing a Matlab .mat file to a *Pandas Dataframe*. Pandas Dataframes hold data in *named* columns with *indexed* rows. Useful tool, numerous use cases.
- Ex. 1: introduces NumPy and a simple vectorization solution.
- Ex. 2: introduces Matplotlib, a few common NumPy “tricks”.

Session 2



- Ex. 0: basic script importing a Matlab .mat file to a *Pandas Dataframe*. Pandas Dataframes hold data in *named* columns with *indexed* rows. Useful tool, numerous use cases.
- Ex. 1: introduces NumPy and a simple vectorization solution.
- Ex. 2: introduces Matplotlib, a few common NumPy “tricks”.

Take-home message:



Session 3



- Ex 3: PSD estimation on non-regularly sampled data.
Main idea: demonstrate SciPy's signal-processing capabilities.

Session 3 (cont')



- Ex 3: PSD estimation on non-regularly sampled data.
Main idea: demonstrate SciPy's signal-processing capabilities.

Take-home message:



Session 3 (cont')



- Ex 4: Cython optimization, profiling.
 - Introduced lprun, Python's line profiler.
 - Introduced Cython, Python's static compiler.
 - Use NumPy arrays as "memoryviews".
 - Demonstrate >1000 (~) speedup over native Python for loops.
 - Demonstrate bounds-check disabling.
 - Demonstrate parallel loop execution.
 - **Take-home messages:**
 - Typically not too difficult to obtain "native" machine speed.
 - Compiling also means worrying about cross-platform, etc.
 - Sometimes the best approach is **brute force**.



Session 4



- Ex 5: Regressions
 - OLS (ordinary least squares) linear regression,
 - Robust linear regression,
 - Ordinary / robust nonlinear (polynomially-inflated) regression,
 - Leave-one-out testing.
 - **Take-home message:** not all statistics are sufficient statistics.

Session 4 (cont')



- Ex 6: feature extraction, parameter tuning / optimization
 - Time series feature extraction (via Pandas' rolling() function),
 - Introduction to Scikit-Learn's API: fit(), predict().
 - Controller parameter optimization via **Optuna's** black-box CMA-ES stochastic optimization.
 - **Take-home messages:**
 - When optimizing over a few ($\leq O(100)$) parameters, **black-box** optimization can easily provide a solution (possibly global, but no performance guarantees).
 - The Optuna library is easy-to-use for such tasks.

Session 5



- Ex 7: distribution drift (a.k.a. *domain shift*, *concept drift*)
 - Demonstration of the (in)famous XOR problem using a linear model.
 - Introduction and use of a few “classical” algorithms:
 - K-nearest-neighbors,
 - Logistic regression.
 - Introduction to a number of useful Scikit-Learn utility functions:
 - `make_pipeline()`,
 - `train_test_split()`.
 - Introduction to Scikit-Learn’s `score()` API.
 - **Take-home message**: beware domain drift.



Session 6 (Midterm Hackathon)



- Ex 8: midterm hackathon
- Take-home messages:
 - Using “all possible” features probably not a good idea.
 - Averaging (over some “natural” timeframe) can de-noise the features, assisting the downstream classification task.
 - If at all possible, decomposing the problem into “natural” sub-problems is typically the best alternative.



Session 7



- Ex 9: Noisy Features

- Toy problem: XOR regression with noisy features.
- Introduced a number of Scikit-Learn regression algorithms:
 - KNeighborsRegressor,
 - DecisionTreeRegressor,
 - AdaBoostRegressor.
- Introduced the `tqdm` class.
- Open problem: when using a 2nd degree polynomial LS regression, why on Earth does the cross-validated loss **spike** when the num. of noise features **equals 4**?
- Clue (?): when using a 3rd degree polynomial LS regression, the loss spikes when the number of noise features **equals 9**.

Session 8



- Ex 10: CatBoost
 - Introduced CatBoostRegressor(), the CatBoost regression API.
 - Provides fast and reliable classification / regression.
 - Does not require feature scaling / normalization / etc.
 - Built-in support for categorical data (not cats).
 - Built-in support for regression uncertainty estimation via the 'RMSEWithUncertainty' loss function.

$$p(y|\mathbf{x}, \boldsymbol{\theta}^{(t)}) = \mathcal{N}(y|\mu^{(t)}, \sigma^{(t)}), \quad \{\mu^{(t)}, \log \sigma^{(t)}\} = F^{(t)}(\mathbf{x}). \quad (9)$$

$$\mathcal{L}(\boldsymbol{\theta}|\mathcal{D}) = \mathbb{E}_{\mathcal{D}}[-\log p(y|\mathbf{x}, \boldsymbol{\theta})] = -\frac{1}{N} \sum_{i=1}^N \log p(y^{(i)}|\mathbf{x}^{(i)}, \boldsymbol{\theta}). \quad (10)$$

Session 8 (cont')



- Take-home messages:
 - CatBoost is a solid baseline when prototyping a new classifier / regressor.
 - Reliable, fast and accurate out-of-the-box performance.
 - Clear advantage over linear regression demonstrated in exercise.
- Additional topics:
 - Parameter tuning (Optuna).
 - Handling class imbalance: `auto_class_weights`.

Session 9



- Ex 11: Dimensionality Reduction
 - Introduced PCA (via Scikit-Learn) and UMAP dimensionality reduction.
 - PCA optimal in capturing maximal variance for given output dimensionality.
 - **Take-home messages:**
 - Proper tool for communication / compression.
 - Not necessarily optimal for classification.
 - Stable ...
 - UMAP optimally embeds a neighbor graph (calculated in the original, high dimensional space) into a low dimensional space.



Session 9 (cont')

- UMAP:
 - Provides high-quality 2D / 3D visualizations,
 - Good starting point for clustering / outlier detection.
 - BUT:
 - Unstable: maps i.i.d. datasets to vastly different lower dim. embeddings.
 - Alternative loss function(s) may be preferable (see PACMAP).
- Additional topics:
 - Outlier Detection Algorithms: `IsolationForest, LocalOutlierFactor.`

Session 10



- Ex 12: Time Series Classification / Regression
- Main tool introduced: MiniROCKET (univariate, multivariate).
- Take-home messages:
 - Quick and easy baseline for time series classification / regression,
 - Achieves $O(\sim \text{optimal})$ performance on analytically tractable toy data,
 - On-par performance on relevant signal processing tasks:
 - Demonstrated (single sine wave) frequency estimation,
 - Demonstrated DTOA (delay estimation) between two noisy, time-shifted signals.
 - When tested on a wide range of real-world (univariate / multivariate) T.S. classification tasks, performance was competitive with (MUCH heavier) state-of-the-art machinery.
 - Begs the hypothesis: real-life (univariate / multivariate) time series probably NOT as complex as 2D (and beyond).

Session 11



- Ex 13: Clustering
 - Introduced HDBScan, a state-of-the-art clustering algorithm,
 - Demonstrated clustering via HDBScan for radar chirp periodicity detection (a.k.a. chaining),
- **Take-home messages:**
 - Provides a quick and solid baseline,
 - Most use cases involve optimizing a single parameter (`min_cluster_size`),
 - Can (in general) do well on up to ~50 to 100 dimensions,
 - Library includes a built-in outlier detection algorithm (GLOSH).

Session 12



- Ex 14: Ensemble methods via CatBoost
- Idea: CatBoost Ensembles (via Posterior Sampling) can provide total uncertainty estimates, as well as a decomposition into:
 - Data (aleatoric) uncertainty,
 - Knowledge (epistemic) uncertainty.
- Methods are considered state-of-the-art in modern NNs.
- **Take-home messages:**
 - Results on toy problem: less than spectacular.
 - Fancy formulas and colorful toy examples should always be taken with a grain of salt.

Session 12



Plain GB:

$$F^{(t)}(\mathbf{x}) = F^{(t-1)}(\mathbf{x}) + \epsilon h^{(t)}(\mathbf{x}), \quad (7)$$

where $F^{(t-1)}$ is a model constructed at the previous iteration, $h^{(t)}(\mathbf{x}) \in \mathcal{H}$ is a weak learner chosen from some family of functions \mathcal{H} , and ϵ is learning rate. The weak learner $h^{(t)}$ is usually chosen to approximate the negative gradient $-g^{(t)}(\mathbf{x}, y) := -\frac{\partial L(y, s)}{\partial s} \Big|_{s=F^{(t-1)}(\mathbf{x})}$:

$$h^{(t)} = \arg \min_{h \in \mathcal{H}} \mathbb{E}_{\mathcal{D}} \left[\left(-g^{(t)}(\mathbf{x}, y) - h(\mathbf{x}) \right)^2 \right]. \quad (8)$$

Stochastic Gradient Langevin Boosting:

$$h^{(t)} = \arg \min_{h \in \mathcal{H}} \mathbb{E}_{\mathcal{D}} \left[\left(-g^{(t)}(\mathbf{x}, y) - h(\mathbf{x}, \phi) + \nu \right)^2 \right], \nu \sim \mathcal{N} \left(0, \frac{2}{\beta \epsilon} I_{|\mathcal{D}|} \right), \quad (11)$$

where β is the inverse diffusion temperature and $I_{|\mathcal{D}|}$ is an identity matrix. This random noise ν helps to explore the solution space in order to find the global optimum and the diffusion temperature controls the level of exploration. Second, the update (7) is modified as:

$$F^{(t)}(\mathbf{x}) = (1 - \gamma \epsilon) F^{(t-1)}(\mathbf{x}) + \epsilon h^{(t)}(\mathbf{x}, \phi^{(t)}), \quad (12)$$



Relegated to Future

- Class imbalance.
- Outlier detection
- Out-of-Distribution (OoD) detection
- Uncertainty estimation
- Classifier calibration:
 - Logistic regression calibrated by-design.
 - Simple NNs already well-calibrated (but modern NNs extremely uncalibrated).
 - Not quite so useful when using CatBoost:
 - Plentiful datasets already fairly well calibrated,
 - Sparse datasets actually hopeless...
 - Calibration useful in various situations not encountered in the exercises (Naive Bayes classifiers, etc.)
- Interpretable ML: SHAP values